

Kaggle Competition Reports

Team: Shivek Ranjan

Shivek Ranjan Roll No: IMT2023042

Navish Roll No: IMT2023060

Digvijaysinh Pawar Roll No: IMT2023099

GitHub Repository: <https://github.com/ImNotFound7/ML-Project-Part-2.git>

November 2025

Contents

I COPD Risk Classification Challenge - Final Report	3
1 Introduction	3
2 What We Learned From Past Attempts	3
2.1 Initial Approach (F1: 0.689)	3
2.2 Key Improvements Made	3
3 Data & Preprocessing	4
3.1 Dataset	4
3.2 Pipeline	4
4 Models & Results	4
4.1 Cross-Validation Performance (10-Fold)	4
4.2 Individual Model Insights	5
5 Stacking Ensemble	5
6 Results Summary	6
7 Breakdown of Improvements	6
8 Key Lessons Learned	6
9 Submission Files	7
10 Conclusion	7
II Signal Cluster Classification Challenge	8

11 Introduction	8
12 Data Description	8
12.1 Dataset Characteristics	9
13 Methodology	10
13.1 Feature Engineering	10
13.2 Modeling	10
14 Results	10
14.1 Cross-Validation Performance	10
14.2 Kaggle Public Leaderboard Performance	11
15 Final Model Selection	11
16 Discussion and Key Learnings	11
17 Conclusion	11
18 Appendix	12

Part I

COPD Risk Classification Challenge

- Final Report

1 Introduction

The COPD Risk Classification Challenge required predicting patient COPD risk (binary classification) from 20 clinical features. Initial baseline: F1 = 0.689 on test set. Target: F1 0.75. Dataset: 44,553 training samples (36.7% COPD positive), 11,139 test samples.

2 What We Learned From Past Attempts

2.1 Initial Approach (F1: 0.689)

Baseline implementation of 5 required models achieved:

- SVM: F1 = 0.70-0.71 (best individual)
- Logistic Regression: F1 = 0.68-0.70
- Neural Network: F1 = 0.62-0.68
- K-Means, GMM: F1 = 0.50-0.65 (unsuitable for supervised task)

2.2 Key Improvements Made

1. Feature Engineering ($20 \rightarrow 68$ features, +2-3% F1):

- Ratios: BMI, cholesterol ratios (LDL/HDL), triglyceride/HDL
- Interactions: age \times BMI, glucose \times BMI, BP \times age
- Risk scores: metabolic syndrome, cardiovascular risk, liver inflammation
- Binary flags: obesity, hypertension, diabetes, dyslipidemia

2. Class Imbalance Handling: SMOTE (+1-2% F1):

- Applied SMOTE before scaling (prevent data leakage)
- Balanced class distribution 50-50 per fold
- Improved recall from 0.67 \rightarrow 0.75

3. Threshold Optimization (+1-3% F1):

- Default threshold (0.5) suboptimal
- Fine-grained search: LR = 0.495, SVM = 0.320, MLP = 0.145, Ensemble = 0.470
- SVM with threshold=0.320 prioritizes catching COPD cases

4. Hyperparameter Tuning (+0.5-1.5% F1):

- LR: C=0.5, penalty='l2', class_weight='balanced'
- SVM: C=1.5, gamma='scale', kernel='rbf', class_weight='balanced'
- MLP: hidden_layers=(256,128,64), alpha=0.0003, learning_rate='adaptive'

5. Stacking Ensemble (+1.25% F1):

- Combined LR + SVM + MLP predictions
- Meta-learner (LR) learned optimal weights
- Outperformed all individual models

3 Data & Preprocessing

3.1 Dataset

- Training: 44,553 samples, 20 features, 36.7% positive class
- Features: age, BMI, BP, glucose, lipids, hemoglobin, liver/kidney enzymes, vision, hearing, oral health

3.2 Pipeline

1. Categorical encoding (sex, oral health, tartar)
2. Feature engineering: 20 → 68 features
3. Imputation: median strategy
4. Scaling: StandardScaler (fit on train only)
5. Feature selection: SelectKBest retaining 60 features

4 Models & Results

4.1 Cross-Validation Performance (10-Fold)

Table 1: Model Performance Comparison

Model	CV F1 Score	Optimal Threshold	Kaggle Score
Logistic Regression	0.7154	0.495	0.711
SVM (RBF)	0.7214	0.320	0.722
Neural Network	0.7142	0.145	0.736
Stacking Ensemble	0.7339	0.470	0.744

4.2 Individual Model Insights

Logistic Regression:

- Stable predictions (std=0.0047)
- Interpretable coefficients
- F1 = 0.7154

SVM (RBF Kernel):

- Captures non-linear relationships
- Optimal threshold = 0.320 (aggressive)
- F1 = 0.7214 (best individual after threshold optimization)

Neural Network:

- High variance (std=0.0104)
- Sensitive to hyperparameters
- F1 = 0.7142

5 Stacking Ensemble

Architecture:

- Level 0: LR, SVM, MLP trained on training folds
- Level 1: Logistic Regression meta-learner
- Input to meta: OOF predictions from 3 base models
- Output: Combined probability with optimal threshold = 0.470

Why Stacking Won:

- Complementary strengths: LR (linear signal) + SVM (non-linearity) + MLP (complex patterns)
- Meta-learner learns optimal combination weights
- Achieves F1 = 0.7339 (best overall)

Table 2: Baseline vs. Final Approach

Component	Baseline	Final
Features	20 raw	68 engineered
Class Balance	Ignored	SMOTE applied
Threshold	Default 0.5	Optimized per model
Ensemble	Simple voting	Stacking
CV Folds	5-fold	10-fold
F1 Score	0.6890	0.7339
Improvement		+6.52%

6 Results Summary

7 Breakdown of Improvements

- Feature engineering: +2-3% F1
- SMOTE balancing: +1-2% F1
- Threshold optimization: +1-3% F1
- Hyperparameter tuning: +0.5-1.5% F1
- Stacking ensemble: +1.25% F1
- 10-fold CV + careful pipeline: +0.5-1% F1

Total: **0.689 → 0.7339 (+0.0449 = +6.52%)**

8 Key Lessons Learned

1. **Feature engineering matters most:** Clinical domain knowledge (ratios, risk scores) provided consistent 2-3% gains
2. **Default threshold (0.5) is suboptimal:** Medical classification benefits from threshold tuning; we found optimal values ranging 0.145-0.495
3. **Class imbalance requires careful handling:** SMOTE improved recall without sacrificing precision; critical for medical applications
4. **Ensemble diversity is key:** Simple voting added marginal value, but stacking with learned meta-weights combined complementary strengths
5. **Unsupervised methods insufficient:** K-Means and GMM unsuitable for supervised classification; importance of labeled data
6. **Simplicity first:** Start with interpretable models (LR) before complex ones; avoid premature complexity

9 Submission Files

Four submissions generated:

- `submission_improved_LR.csv`: F1 training = 0.7154, threshold = 0.495
- `submission_improved_SVM.csv`: F1 training = 0.7214, threshold = 0.320
- `submission_improved_MLP.csv`: F1 training = 0.7142, threshold = 0.145
- `submission_improved_STACKED.csv`: F1 training = 0.7339, threshold = 0.470

Also saved: `submission_improved_PROBABILITIES.csv` with raw probabilities from all models for further analysis.

10 Conclusion

Starting from baseline F1 = 0.689, systematic improvements achieved final F1 = 0.7339 (+6.52% improvement) through:

1. Advanced feature engineering capturing clinical syndromes
2. Proper class imbalance handling with SMOTE
3. Fine-grained threshold optimization
4. Effective stacking ensemble combining complementary models
5. Careful hyperparameter tuning

The stacking ensemble with threshold = 0.470 balances precision (0.7210) and recall (0.7470), making it ideal for medical diagnosis where both false negatives and false positives are costly.

Recommendation: Submit `submission_improved_STACKED.csv` for highest F1 score on test set.

Technical Details

Implementation: Single Jupyter notebook (`test-cl.ipynb`) with reproducibility (`random_state=42`)

Requirements: Python 3.11, pandas, numpy, scikit-learn, imbalanced-learn

Code Structure:

1. Data loading & preprocessing
2. Feature engineering (20 → 68)
3. 10-fold stratified CV with SMOTE
4. Per-model training with GridSearch
5. Threshold optimization on OOF
6. Stacking ensemble training
7. Test predictions & submission generation

Part II

Signal Cluster Classification Challenge

11 Introduction

This report documents our participation in the Kaggle competition **Signal Cluster Classification**. The objective is to predict the *personality_cluster* (class label) for a given signal point defined by two continuous features: *signal_strength* and *response_level*.

Initially, the task appeared to be a simple 2D classification problem, but achieving high performance on the hidden leaderboard required careful feature engineering, kernel-based modeling, and extensive validation.

12 Data Description

The dataset consists of two numeric features and one target label:

- Signal Strength
- Response Level
- Personality Cluster (categorical target)

The data represents synthetic clusters in a 2D space. This structure makes the task a nonlinear, geometry-driven multiclass classification problem.




Figure 1: Distribution of Signal Strength and Response Level in Train and Test Data

12.1 Dataset Characteristics

The competition dataset is divided into two separate files: `train.csv` and `test.csv`. The training set contains **1444 samples**, each with two numeric signal features and one target label. The test set contains **362 samples** for which predictions must be generated.

There are **three distinct cluster classes** in the dataset. Class distribution is relatively balanced, making Macro F1 Score an appropriate evaluation metric.

Both features exhibit a wide dynamic range with clear nonlinear distribution. For the training dataset:

- Signal Strength ranges from **-62.58** to **756.04**
- Response Level ranges from **-55.08** to **543.57**

The test set maintains similar characteristics:

- Signal Strength ranges from **-26.02** to **787.82**
- Response Level ranges from **-33.69** to **539.97**

This confirms that the training and test distributions are aligned closely enough to trust cross-validation performance while still requiring robust generalization to slight cluster shifts.

13 Methodology

Our approach evolved through multiple experiment cycles. The workflow consisted of:

1. Data exploration and visualization
2. Feature engineering based on geometry
3. Model benchmarking with strict cross-validation
4. Leaderboard-based generalization testing

13.1 Feature Engineering

We derived expressive features to capture nonlinear spatial patterns:

- Polynomial: x^2, y^2, xy, x^3, y^3
- Radial distance: $r = \sqrt{x^2 + y^2}$ and r^2, r^3
- Angular: $\theta = \arctan 2(y, x)$ and trigonometric expansions
- KMeans cluster distances (for geometry adaptation)

These helped the model approximate complex cluster boundaries.

13.2 Modeling

We trained and evaluated three major models:

- **SVM with RBF Kernel** — best performance
- **MLP Neural Network** — strong but slightly less accurate
- **Stacking Ensemble (SVM → MLP)** — surprisingly worse due to error propagation

Cross-validation was performed using **10-fold Stratified K-Fold** with Macro F1 score.

14 Results

14.1 Cross-Validation Performance

Model	CV Macro F1
MLP Neural Network	0.9858
SVM RBF	0.9833
Stacking Classifier	0.9810

Table 3: Cross-Validation Results

Model	Public F1 Score
SVM RBF	0.989
MLP Neural Network	0.985
Stacking Classifier	0.977

Table 4: Leaderboard Results

14.2 Kaggle Public Leaderboard Performance

Despite the lower CV score, SVM generalized the best to the test set.

This mismatch highlights a small domain shift between the training and test distributions.

15 Final Model Selection

The **SVM with RBF kernel** was selected as the final model based on its superior leaderboard performance. The winning configuration featured:

- Well-engineered geometric features
- Class-weight balancing
- Hyperparameter tuning over a competitive search space

16 Discussion and Key Learnings

1. **Best model in CV is not always best on test:** MLP had higher CV but underperformed on the leaderboard.
2. **Stacking can reduce performance:** Poor error correction made the stacked model weaker.
3. **Feature engineering mattered more than architecture:** Simple 2D input became highly separable through geometry-based transformations.
4. **SVM excels in boundary-driven datasets:** The RBF kernel created smooth, generalizable decision contours.

17 Conclusion

Through structured experimentation, the SVM RBF model emerged as the optimal solution, achieving a strong leaderboard score of **0.989**. Our results reinforce:

“Deep understanding of the data geometry beats blindly adding complexity.”

The accompanied working code and predictions ensure full reproducibility for grading and submission.

18 Appendix

- Requirements: Python 3.11, pandas, numpy, scikit-learn